

# AI runs the work. A human holds every gate.

How Evalyn — an AI-operated IT services company — works: the organization, the value chain, the approval loop, the architecture, the 7 lifecycle workflows, and a one-page dossier for each of the 23 AI employees.

|                                       |                                     |                                      |                                   |                                             |
|---------------------------------------|-------------------------------------|--------------------------------------|-----------------------------------|---------------------------------------------|
| <div>23</div> <div>AI EMPLOYEES</div> | <div>7</div> <div>DEPARTMENTS</div> | <div>26</div> <div>ROLE SKILLS</div> | <div>7</div> <div>WORKFLOWS</div> | <div>100%</div> <div>GATES HUMAN-HELD</div> |
|---------------------------------------|-------------------------------------|--------------------------------------|-----------------------------------|---------------------------------------------|

# Contents

|    |                                                  |                         |
|----|--------------------------------------------------|-------------------------|
| 01 | <b>The organization</b>                          |                         |
| 02 | <b>The value chain — lead to cash to renewal</b> |                         |
| 03 | <b>Cross-team work — inside one stage</b>        |                         |
| 04 | <b>The approval loop</b>                         |                         |
| 05 | <b>The architecture</b>                          |                         |
| 06 | <b>The roadmap</b>                               |                         |
| 07 | <b>The 7 lifecycle workflows</b>                 | — one page each         |
| 08 | <b>The 23 AI employees</b>                       | — one dossier page each |

How to read this document: **teal** marks AI doing work; **amber** marks a point where a named human decides; dashed boxes are records in the Company OS — the shared system of record every employee reads before acting and writes after.

# The organization

Every department has three human seats — **Approver, Deputy, Head** — that AI work escalates to. Today the founder-operator holds every seat (the model works at  $n=1$ ); hiring hands seats over one department at a time.

## Founder / CEO / Human Operator

♦ **TERMINAL BACKSTOP**

End of every escalation chain. No approval or question can ever dead-end.

### ■ LEADERSHIP

ceo eng-manager

gates: strategy · escalations

Seats: **Approver · Deputy · Head — Founder-Operator**

### ■ PEOPLE & FINANCE

hr finance

gates: money · people (dual) · revenue booking

Seats: **Approver · Deputy · Head — Founder-Operator**

### ■ PRODUCT & DESIGN

product-manager project-manager  
designer tech-writer

gates: roadmap commitments

Seats: **Approver · Deputy · Head — Founder-Operator**

### ■ ENGINEERING

developer code-reviewer tester  
devops security

gates: merge · deploy · migrations

Seats: **Approver · Deputy · Head — Founder-Operator**

### ■ SALES & DELIVERY

sdr sales solutions-architect  
delivery-manager account-manager

gates: price · date · SLA commitments

Seats: **Approver · Deputy · Head — Founder-Operator**

### ■ MARKETING & SUPPORT

marketing social-media support

gates: external comms · publish

Seats: **Approver · Deputy · Head — Founder-Operator**

### ■ OPERATIONS


procurement data-analyst

gates: orders · vendor signing

Seats: **Approver · Deputy · Head — Founder-Operator**

# The value chain

AI employees hand finished artifacts down the chain; every amber pill is a point where the flow **stops until a human approves**. Reading connected systems is free — anything sent outside the company is gated.



## The seven human gates (CLAUDE.md §5 — absolute)



## Inside one stage: cross-team work

A stage is rarely one employee. Each stage has **one accountable lead** who fans the work out to specialists across departments — in parallel — then reconciles their inputs into the stage's record. The lifecycle workflows are exactly this orchestration, scripted so it runs the same way every time.



# The approval loop

The heart of the minimum-human design. An AI employee never waits in a chat for an answer — it **writes a record and stops**; routing guarantees the record reaches a human who can decide.

1

## AI employee hits a gate — writes a record and stops

The artifact is finished. The employee writes an approval request with the **exact action** — "send proposal v3 to jane@acme.com" — never a vague ask. Questions to humans travel the same way.

```
company/approvals/APR-20260712-001.md
```

```
company/questions/QST-....md
```

2

## Routing finds an available human — and escalates if none responds

The gate type maps to a department; the record is assigned to its Approver. Unavailable humans are skipped instantly; an expired SLA escalates one hop. Every hop is logged.

**Dept Approver**

 unavailable /  
SLA breach →

**Deputy**
**Dept Head**
**CEO — always staffed**
**P0** urgent — 2 h, escalate hourly

**P1** normal — 1 business day

**P2** low — 3 business days

3

## A human decides — in the Control Panel inbox

The inbox shows the artifact, the AI's summary and recommendation, and the exact action. Four moves, every one stamped with who / when / why: **Approve** (optional conditions) · **Reject** (reason required, feeds rework) · **Answer** (for questions) · **Delegate** (to another authorized human).

4

## Approved action executes — exactly once, verbatim

The executor claims the record, performs the action exactly as approved (artifact hash checked — anything changed since approval goes back for re-approval), then stamps it executed. A claim that never completes goes to a human — never a silent retry.

```
execution: {claimed_at → executed_at}
```

◆ **Silence never equals consent. An expired SLA escalates to the next human — nothing is ever auto-approved, auto-rejected, or auto-sent. (ADR-0004)**

# The architecture — one git repository is the whole company

The repo is the **single source of truth and the audit log**. Everything else — AI sessions, the SLA job, the Control Panel — is a stateless client reading it and writing commits. If the panel is down, the file-based loop still works.

## AI employees CLAUDE

### CODE

23 personas run as agent sessions. They do the work, write Company OS records, and stop at gates.

writes → records + approval requests

## SLA / escalation job CRON

Every ~15 min: scan pending items, apply routing, escalate breaches, send email notifications. Idempotent — a missed run catches up.

writes → hops + notifications

## The git repository

SOURCE OF TRUTH · AUDIT LOG · MEMORY

CLAUDE.md — the operating brain

.claude/ — agents · skills · workflows

company/ — the business records

org/ — humans · seats ·

routing · SLAs

approvals/ · questions/

leads/ ... invoices/ ... assets/

docs/ — plans · specs · ADRs

memory/ — cross-session memory

## Control Panel WEB

### HUMANS

Company dashboard, department boards, and the approval inbox. Decisions become commits with the human's identity stamped.

reads state → writes decisions

## Humans APPROVERS

Notified by email (Slack later) with deep links. The panel is a window, never the only path.

approve · reject · answer · delegate

## Why git-native

Git history **is** the audit trail — attribution, immutability, and diffs for free. Near-zero new infrastructure for the first phases. A dedicated backend comes only when latency measurably hurts (Phase 4), and the repo stays authoritative even then.

# Inside the Control Panel — how one decision physically happens

The panel is a **stateless web app over the repo** (Tech Spec 001): it reads records through a short-lived cache and writes every decision back as a commit stamped with the human's identity. No database — git history is the audit log.

## Humans — Approver · Deputy · Head · CEO

Phone or desktop. Identity maps 1:1 to company/org/humans/ — no shared logins; only authorized seats can decide an item.

## Email notifications

Sent on **assigned** · **50% SLA elapsed** · **escalated** — each with a deep link straight to the item. Slack arrives in Phase 3.

↑ notifies with deep link

↓ signs in · decides

## Control Panel

STATELESS WEB APP · SERVER-RENDERED · NO DATABASE

### AUTH & AUTHORIZATION

Email magic-link → session must match a record in **company/org/humans/** → every decision write is authorized server-side against **routing.md** (the right seat, the right department, dual-approval where required).

### Approval Inbox

### Item Detail

### Company Dashboard

### Department Board

### People & Routing Admin

### Availability Toggle

### READ PATH

GitHub API → server-side cache, TTL ≤ 5 min, revalidated after the panel's own writes. Inbox sorted by SLA urgency, worst first.

### WRITE PATH

One decision = **one commit** on the record: stamp {who · when · outcome · why} + Decided-by trailer. Optimistic concurrency — a record changed since read reloads instead of clobbering.

↓ reads state

↓ writes decision commits

## SLA / escalation job

cron ~15 min

Scans state: pending → skips unavailable humans, escalates SLA breaches one hop, appends to the record, sends the emails. Idempotent — a missed run catches up.

## The git repository

company/approvals/ · questions/ · org/ — every record carries its state, assignee, hop history, decision stamp, and execution stamp. Git history **is** the audit trail.

## Executor on approval

Picks up approved + unexecuted → **claims** the record → verifies the artifact hash matches approval time → performs the exact action once → stamps executed\_at.

1

**AI hits a gate** — writes the approval record with the exact action, and stops.

2

**SLA job routes it** — assigns the available approver, emails a deep link.

3

**Human opens the inbox** — artifact, AI summary, exact action.

4

**Decision becomes a commit** — approve / reject / answer / delegate, stamped.

5

**Executor runs it exactly once** — claim → hash check → execute → stamp.

# The roadmap

Evalyn builds this as its own first internal project, through its own value chain — **every merge approved by the very system under construction.**

Week 1

## Phase 0 — Decide, don't build IN PROGRESS

PRD, tech spec, and ADRs written, awaiting sign-off. Then: CLAUDE.md gate model updated, company/org/ created with the founder-operator in every seat, one sample gate dry-run on paper.

Weeks 2–3

## Phase 1 — Prove the loop, zero UI NEXT

File-based approval flow live: records, SLA job, email, exactly-once execution. One real deliverable end to end, plus a deliberate escalation drill.

Weeks 4–6

## Phase 2 — Control Panel MVP NEXT

Inbox, company dashboard, department boards, availability toggle. Exit test: a non-technical human approves an AI delivery without touching git.

Weeks 7+

## Phase 3 — Automation & first hires NEXT

Auto-reassignment, delegation UI, Slack, audit screen — and the first real Department Approvers take seats, one department at a time.

When it hurts

## Phases 4-5 — Scale-up & authoring skills LATER

Dedicated backend only when git latency measurably hurts. Gated /create-agent · /create-skill · /create-workflow when client work demands roster growth.


## Success metrics

Median time-to-decision < SLA/2 · zero terminal-SLA breaches · zero dead-ended items · 100% of gated actions traceable to a named human · an approver handles an item in < 5 minutes median.

# opportunity-to-proposal

Move a qualified opportunity to a sendable proposal: discovery → scope → estimate → pricing → proposal, with an internal deal-risk review before anything reaches the customer.

**Use when:** a qualified opportunity needs a proposal. Input: the opportunity id, what the customer wants, and the commercial model (fixed / T&M / retainer).



# project-kickoff

Turn a booked PO into a ready-to-deliver project: verify the PO against the deal, then build the SOW, resourcing, and delivery plan in parallel, then the risk register and kickoff pack.

**Use when:** a customer PO is in and the project must be set up for delivery. Input: project id, PO ref + amount + terms, link to the quote/proposal.

## 1 · Verify

Reconcile PO to the deal — catch mismatches before they become disputes

→

## 2 · Plan

SOW · resourcing · delivery plan & milestones

|| PARALLEL

→

## 3 · Risk

Delivery risk register

→

## 4 · Kickoff

Assemble the kickoff pack

### EMPLOYEES INVOLVED

delivery-manager

solutions-architect

eng-manager

project-manager

finance

### HUMAN GATES

**Committing dates & resourcing to the client**

### READS / WRITES

pos/ + proposals/ → sows/ → projects/ → milestones/


### OUTPUT

A kickoff pack the eng pod can execute from day one: SOW, milestone schedule with acceptance criteria, named risks with mitigations.

# delivery-to-invoice

Close the delivery→cash loop for a milestone: verify the deliverable against acceptance criteria, run the QA/security gate, prepare client acceptance, draft the invoice reconciled to the PO/SOW.

**Use when:** a milestone is delivered and should be accepted and billed. Input: project id, milestone, acceptance criteria from the SOW.



EMPLOYEES INVOLVED

delivery-manager

tester

security

finance

HUMAN GATES

Acceptance sign-off · sending the invoice · collections actions

READS / WRITES

milestones/ + sows/ + pos/ → invoices/

WHY IT MATTERS

Acceptance triggers billing — so the invoice is reconciled to the PO/SOW before a human ever sees it. No unbacked line items.

# procurement-cycle

Run a buying need to a ready-to-approve order: check the asset register for reclaim first, source and compare vendors on total cost of ownership, risk-check, draft the outbound PO.

**Use when:** the company needs to buy something — tools, licenses, cloud, hardware, subcontractor. Input: what + why + for whom, quantity, budget cap.

## 1 · Reclaim-check

Can existing assets cover this? Cheapest purchase is none.

→

## 2 · Source

Vendor options compared on TCO  
|| PARALLEL

→

## 3 · Risk

Security / data risk + TCO review

→

## 4 · Order

Recommendation + draft PO

### EMPLOYEES INVOLVED

procurement

security

finance

### HUMAN GATES

**Approving spend · placing the order · signing the vendor**

### READS / WRITES

assets/ → vendors/ → purchase-orders-out/

### THEN

On approval, the asset is registered with owner, renewal date, and utilization tracking — waste surfaces for reclaim next cycle.

# hiring-pipeline

Run the hiring prep for a role: success profile, then JD, sourcing strategy, structured interview loop, and bias-aware scorecard built in parallel, assembled into a fair hiring kit.

**Use when:** opening a role — including Evalyn's own future Department Approvers. Input: role title + level, stack, location, budget, must-haves.

## 1 · Profile

Success profile + competencies

→

## 2 · Build

JD · sourcing · interview loop ·  
scorecard

|| PARALLEL

→

## 3 · Assemble

The hiring kit + fairness check

### EMPLOYEES INVOLVED

**hr** **eng-manager** (+ the hiring department's lead)

### HUMAN GATES

**The hire decision · offers · compensation —  
always human**

### OUTPUT

A complete kit: JD, sourcing plan, structured loop with question banks, bias-aware scorecard, and a fair debrief framework.

### FAIRNESS BY CONSTRUCTION

Same structured process for every candidate;  
scorecards anchored to evidence, not impressions.

# product-launch

Coordinate a cross-functional launch: every function prepares its launch readiness in parallel, then one launch-readiness verdict with blockers and a go/no-go recommendation.

**Use when:** launching a feature or product across the whole company. Input: what is launching, target date, scope (beta / GA / major).

## 1 · Prep

Each function prepares readiness: product · engineering/QA · docs · marketing · sales · support · social

|| 7 FUNCTIONS IN PARALLEL

→

## 2 · Verdict

Launch-readiness go / no-go with blockers

### EMPLOYEES INVOLVED

product-manager   eng-manager   tester  
tech-writer   marketing   sales   support  
social-media

### HUMAN GATES

**Every public action — the announcement, the posts, the pricing page**

### OUTPUT

A readiness matrix per function, the blocker list with owners, and a go/no-go a human can act on.

### PRINCIPLE

No function's "ready" is assumed — each proves it in parallel, and one missing function blocks the verdict, not the discovery of it post-launch.

# company-standup

Roll up a company-wide status: each department reports in parallel, then a CEO-level synthesis with risks and the one decision needed. The at-a-glance heartbeat of the whole company.

**Use when:** a weekly (or on-demand) company status. Input: period ("this week") and an optional focus area. Also the reporting cadence for the Control Panel build itself.

## 1 · Departments

Engineering · product · delivery · GTM · people/finance · support/data each report status, risks, asks

|| 6 DEPARTMENT REPORTS IN PARALLEL

→

## 2 · Synthesize

CEO-level roll-up: on track / at risk / blocked + the one decision needed

### EMPLOYEES INVOLVED



### HUMAN GATES

**None — internal reporting. Decisions it surfaces route to their own gates.**

### OUTPUT

One page: what moved, what's at risk (with mitigation), what's blocked (with the ask), and the single decision leadership must make.

### FEEDS

memory/decisions-log.md · the Control Panel company dashboard


# The 23 AI employees

Each dossier page reads the same way: what the employee **receives**, what it **owns & does**, what it **hands off** and to whom — then its skills, its amber line (what it never does without a human), and when it escalates. Together they are the org: every handoff below has a receiver, every gate has an approver.

|                                                      |                                                |                                                      |
|------------------------------------------------------|------------------------------------------------|------------------------------------------------------|
| <b>ceo</b><br>Chief Executive 01                     | <b>eng-manager</b><br>Engineering Manager 02   | <b>hr</b><br>People / HR Partner 03                  |
| <b>finance</b><br>Finance Partner 04                 | <b>product-manager</b><br>Product Manager 05   | <b>project-manager</b><br>Project Manager (Scrum) 06 |
| <b>designer</b><br>Product Designer 07               | <b>tech-writer</b><br>Technical Writer 08      | <b>developer</b><br>Software Engineer 09             |
| <b>code-reviewer</b><br>Code Reviewer 10             | <b>tester</b><br>QA Engineer 11                | <b>devops</b><br>DevOps / SRE 12                     |
| <b>security</b><br>Security Engineer 13              | <b>sdr</b><br>Sales Development Rep 14         | <b>sales</b><br>Account Executive 15                 |
| <b>solutions-architect</b><br>Solutions Architect 16 | <b>delivery-manager</b><br>Delivery Manager 17 | <b>account-manager</b><br>Account Manager / CS 18    |
| <b>marketing</b><br>Marketing Lead 19                | <b>social-media</b><br>Social Media Manager 20 | <b>support</b><br>Support Engineer 21                |
| <b>procurement</b><br>Procurement & Assets 22        | <b>data-analyst</b><br>Data Analyst 23         |                                                      |

## CEO AI Chief Executive

Sets strategy and priorities, allocates focus, resolves cross-functional conflicts, and prepares board/investor communications. Thinks in outcomes and trade-offs, not tasks.



### SKILLS

/okrs

/board-update

/strategy-review

### NEVER WITHOUT A HUMAN


**Contracts & public statements · financing actions · headcount / compensation · spend approval**

### ESCALATES WHEN

A decision is irreversible and high-stakes; two departments deadlock strategically; a material legal/financial/reputational/security risk appears. Always options + a recommendation, never just the problem.

## eng-manager AI Engineering Manager

Owns engineering delivery and team health: unblocks the pod, balances load, reports status up, and holds the process gates before merge and deploy.



### SKILLS

/one-on-one

/team-report

/sprint-plan

code-review

deploy-checklist

### NEVER WITHOUT A HUMAN


**Approving risky merges / prod deploys · committing delivery dates to a customer · any people decision (rating, PIP, role change)**

### ESCALATES WHEN

A commitment is at risk and unrecoverable inside the team; priorities conflict beyond authority; a people/security/reliability risk is material.

## hr AI People / HR Partner

Owns hiring, onboarding, performance process, culture, and policy. Structures every people decision fairly — and never makes the hire/fire/comp call itself.



### SKILLS

/job-description

/hiring-plan

/onboarding-plan

/performance-review

### NEVER WITHOUT A HUMAN


**Offers & rejections · terminations / PIPs · ratings & compensation · any candidate or employee communication**

### ESCALATES WHEN

Legal/compliance risk (discrimination, wrongful-termination exposure); a request conflicts with fair process; any sensitive people situation — routed to a human, never resolved autonomously.

## finance AI Finance Partner

Owns budgeting, runway/forecasting, financial modeling, and invoicing prep. Computes with real numbers — never estimates a figure it can calculate; unknowns are TBD.



### SKILLS

/budget-plan

/financial-model

/valuation

/invoice

### NEVER WITHOUT A HUMAN


**Moving money · approving spend · sending an invoice · committing financial terms · sharing financials externally**

### ESCALATES WHEN

Runway drops below a safe threshold; a scenario shows obligations can't be met; numbers don't reconcile; a spend request exceeds policy. The number, the cause, the options.

## product-manager AI Product Manager

Decides what to build and why. Owns specs, roadmap, and prioritization grounded in user problems and data — and connects shipped work back to outcomes.



### SKILLS

/spec

roadmap

/metrics-review

### NEVER WITHOUT A HUMAN


**Public or customer feature/date commitments · roadmap promises to the market · deprecating anything customers rely on**

### ESCALATES WHEN

Priorities conflict beyond authority; evidence contradicts a leadership assumption; a commitment is requested the team can't safely meet.

## project-manager AI Project / Delivery Manager (Scrum)

Owns delivery mechanics: sprint planning, standups, tracking, dependencies, and risk. Coordinates — does not decide what to build or how.



### SKILLS

/sprint-plan

/standup

risk registers

### NEVER WITHOUT A HUMAN

**Committing a delivery date externally · scope changes that affect a customer commitment**


### ESCALATES WHEN

The sprint goal is unrecoverable inside the team; a dependency will slip a commitment; scope and capacity are irreconcilable. Always with options.

# designer

## AI Product Designer

Owns UX/UI, user flows, the design system, and prototypes. Turns a product problem into a usable, accessible experience — every state designed, nothing left implicit.



### SKILLS

/design-brief

design-critique

accessibility-review

design-handoff

ux-copy

### NEVER WITHOUT A HUMAN


**Treating a design as final-for-build without sign-off · designs that imply a product/scope change**

### ESCALATES WHEN

The requested UX conflicts with usability/accessibility; the problem is under-defined to design against; a decision has product implications beyond the lane.

## tech-writer AI Technical Writer

Owns documentation: user guides, API docs, READMEs, runbooks, release notes. Grounds every doc in the actual product; marks unknowns TBD, never invents behavior.



### SKILLS

/documentation

release-notes

### NEVER WITHOUT A HUMAN

**Publishing anything externally · disclosing unreleased features or security-sensitive detail**


### ESCALATES WHEN

Product behavior contradicts the intended docs (possible bug); a feature is too confusing to document cleanly (a design issue); a critical detail can't be confirmed.

# developer

## AI Software Engineer

Implements features, fixes bugs, writes tests, and opens PRs. Production-quality, tested, reviewable code — and a full stop at the merge/deploy gate.



### SKILLS

/refactor-plan

test-writer

adversarial-verifier

code-review

### NEVER WITHOUT A HUMAN


**Merging to main · deploying · database migrations · touching production data · risky dependencies · large-scale deletions**

### ESCALATES WHEN

The spec is ambiguous; the blast radius grows beyond expectation; the decision is above the lane (architecture change, breaking API). A specific question with options — to product-manager / eng-manager.

## code-reviewer AI Code Reviewer

Reviews changes for correctness, security, performance, and maintainability before merge. Verified, ranked findings only — approval to merge stays with a human.



### SKILLS

code-review

adversarial-verifier

### NEVER WITHOUT A HUMAN


**Merge approval is a human decision — recommendation only (approve / approve-with-nits / request-changes). Never signals "merged."**

### ESCALATES WHEN

A critical security or data-integrity issue appears (loop in security); the change needs an architectural decision; the diff is too large to review safely — ask to split it.

## tester AI QA Engineer

Owns test strategy, test generation, regression, and bug triage. Verifies quality before release with evidence — and gives a go/no-go a human can act on.



### SKILLS

/test-strategy test-gap test-writer /ticket-triage

### NEVER WITHOUT A HUMAN


**Release sign-off is advisory: ready / ready-with-risks / not-ready, with evidence. Never marks anything shipped.**

### ESCALATES WHEN

A release-blocking defect is found; coverage is too thin to sign off safely; a bug reveals a systemic issue — bad spec, missing validation.

## devops AI DevOps / SRE Engineer

Owns CI/CD, deploys, infrastructure, and incident response. Produces the go/no-go and the rollback plan — a human authorizes every production change.



### SKILLS

deploy-readiness

/incident

/postmortem

IaC

### NEVER WITHOUT A HUMAN

**Production deploys · migrations · prod infra changes · secrets rotation · any destructive operation**

### ESCALATES WHEN

Readiness is NO-GO; an incident is SEV1; a rollback isn't clean; an infra change carries data-loss risk. Verdict, impact, options.

## security AI Security Engineer (AppSec)

Owns security review, threat modeling, and vulnerability triage — including LLM/ML security for AI features. Verifies exploitability; never ships fixes itself.

### RECEIVES

- Diffs & designs to review
- Reported vulnerabilities
- Vendor-risk checks from **procurement**

→

### OWNS & DOES

- Security review of code & design
- Threat models per feature
- Vuln triage with severity & evidence
- Secure defaults: authn/z, secrets, validation

→

### HANDS OFF

- Fixes → **developer**
- Ship-blockers → **eng-manager / devops** + human
- Breach signals → human immediately

### SKILLS

code-review (security lens)

ml-security-audit

red-teamer

### NEVER WITHOUT A HUMAN


**Implementing or deploying fixes itself · disclosing vulnerabilities externally**

### ESCALATES WHEN

A critical exploitable vulnerability exists (data exposure, RCE); a fix is being skipped; an incident looks like a breach. Proof, impact, remediation.

## sdr AI Sales Development Rep

Owns the top of the funnel: generates and qualifies leads against the ICP, drafts first outreach, and converts qualified leads into opportunities in the system of record.



### SKILLS

/lead-gen

/qualify-lead

/sales-outreach

deep-research

### NEVER WITHOUT A HUMAN


**Sending any external message — drafts are queued; a human sends. Internal lead/opportunity records are free to write.**

### ESCALATES WHEN

A lead can't be qualified without technical depth (→ solutions-architect), or the ICP itself seems wrong for the market it's finding (→ ceo/sales).

# sales AI Sales / Account Executive

Owns outreach, discovery, proposals, and the pipeline. Researched, honest, personalized drafts — a human sends every external message and signs everything.



SKILLS

- /sales-outreach
- /proposal
- discovery-synthesis

**NEVER WITHOUT A HUMAN**


**Sending external messages · quoting price / discount / terms · committing dates or SLAs · signing or agreeing to anything**

**ESCALATES WHEN**

A deal needs a non-standard term or discount; a prospect requests commitments beyond the roadmap; a legal/contractual question arises.

## solutions-architect AI Solutions Architect (pre-sales)

Owns discovery, technical scoping, and effort estimation for prospective projects. Turns a customer need into a scoped solution and a defensible estimate — flags feasibility risks honestly, never inflates to win.



### SKILLS

/estimate

/adr

system-design

deep-research

### NEVER WITHOUT A HUMAN


**Sharing scope or estimates with the customer — only via the human-approved proposal**

### ESCALATES WHEN

Feasibility is doubtful or the ask exceeds capability — flag to sales/eng-manager with the risk. The honesty bar beats winning the deal.

## delivery-manager AI Delivery Manager

Owns client project delivery end to end after a deal is won: SOW, resourcing, kickoff, milestones, acceptance, and client status. The link between delivery and billing.



### SKILLS

[/sow](#)[/kickoff](#)[/sprint-plan](#)

### NEVER WITHOUT A HUMAN


**Committing dates/resourcing to the client · agreeing change orders · marking a milestone accepted (it triggers invoicing)**

### ESCALATES WHEN

Delivery is at risk against a commitment → eng-manager/ceo with options. Never hides a slip from the client status.

## account-manager AI Account Manager / Customer Success

Owns the customer relationship after go-live: health, QBRs, renewals, and honest upsell grounded in the customer's real usage — and advocates for the customer internally.



### SKILLS

/qbr

/proposal

/metrics-review

### NEVER WITHOUT A HUMAN


**Sending anything to the customer · committing pricing/terms · retention concessions**

### ESCALATES WHEN

Churn risk is material — the signal plus a save plan, early. A delivery or support issue threatens the account → loop the owners + ceo.

## marketing AI Marketing Lead

Owns campaigns, content, SEO, positioning, and launches. Every claim backable, every campaign with one goal and one measurable metric. Drafts and stages; a human publishes.



### SKILLS

/campaign-brief

/content-calendar

/seo-brief

### NEVER WITHOUT A HUMAN

**Publishing anything public · committing marketing spend · public product or roadmap claims**

### ESCALATES WHEN

A claim can't be substantiated (→ product-manager); a campaign needs budget beyond policy (→ finance); messaging conflicts with product or legal.

## social-media AI Social Media Manager

Owns the social calendar, post drafting, and community engagement — each platform in its own voice. Drafts everything; a human approves before anything goes public.

### RECEIVES

- Calendar & campaigns from **marketing**
- Launch plans
- Mentions & community signals

→

### OWNS & DOES

- The social content calendar
- Platform-tailored post drafts + variants
- Engagement/response drafts
- Capturing what resonates

→

### HANDS OFF

- Strategy alignment → **marketing**
- Performance → **data-analyst**
- Escalations → **ceo** + human

### SKILLS

/social-post

/content-calendar

### NEVER WITHOUT A HUMAN


**Publishing any post, reply, or DM · public responses to complaints/criticism · anything PR-sensitive**

### ESCALATES WHEN

A post could be a PR risk; a complaint escalates publicly; a trend demands a fast company response — never freelances a crisis response.

## support AI Customer Support Engineer

Owns ticket triage, bug reproduction, and customer response drafting — empathetic, policy-grounded replies, and clean repro handed to engineering. A human sends every reply.



### SKILLS

/ticket-triage

/support-macro

### NEVER WITHOUT A HUMAN


**Sending any external customer message · promising a fix, refund, or timeline · committing anything the company must honor**

### ESCALATES WHEN

A ticket is a security/privacy report; a customer is at churn or legal risk; a bug is widespread — possible incident, straight to devops.

## procurement AI Procurement & Asset Manager

Owens buying (vendors, quotes, outbound POs) and inventory (licenses, devices, cloud). Checks reclaim before buying, compares on TCO, and never commits spend itself.



### SKILLS

/procurement-request

/asset-register

deep-research

### NEVER WITHOUT A HUMAN

**Placing any order · committing spend · signing a vendor agreement**


### ESCALATES WHEN

Spend exceeds policy; a vendor poses a security/data risk; a needed asset is blocking delivery. Waste (unused/duplicate assets) surfaced for reclaim.

# data-analyst

## AI Data Analyst

Turns product and business data into trustworthy metrics, dashboards, and insight. Computes every number it reports — never guesses, never exposes PII.



### SKILLS

[/metrics-review](#)[dataviz](#)[data-validator](#)

### NEVER WITHOUT A HUMAN

**Analysis is advisory — humans decide · sharing sensitive metrics externally · any PII in outputs**

### ESCALATES WHEN

Data quality is too poor to trust; a metric reveals a serious problem (churn spike, revenue drop, SLA breach); numbers can't be reconciled.